

KV cache compression

1) Quantisation

Same as for Model comp.

2) Eviction Policies :

drop KV caches for tokens that matter the least important.

How it works :

1) Heavy-hitter strat.

tokens that historically get high score stay while those which consistently get ignored are evicted

↳ minimal quality loss

2) Sliding window + Heavy hitter :

keep most recent w tokens +

top k heavy hitters from older

3) Grouped Query Attention / GQA

- Multi Query attention

while doing multi head attention instead of each head having its own Q, K, V values, we create groups with similar goals & make them share same K, V . Still have own Q .

GQA $\rightarrow$ groups share same K, V

MQA $\rightarrow$ all heads share same K, V

